# Kuku AI Generalist Assignment — Problem 1: Find the Promos

Version 2 | Dialogue-aware, event-aligned screenplay promo candidate finder

## Problem and approach

Editors need to identify short promotional moments from episodic scripts. This prototype parses timestamped screenplay beats, proposes 25–60 second windows and ranks them for human review. It processes the supplied screenplay ZIP; it does not edit video or claim to predict ad performance.

## What changed in Version 2

The parser uses screenplay section labels to distinguish dialogue, action and camera beats. Candidate starts align to timed beat boundaries. The scorer rewards dialogue evidence near the opening and ending, counts tension-related terms, considers beat density and uses the episode logline as lightweight context. Candidates with substantial interval overlap are suppressed. Visual/action fallbacks are labelled rather than presented as dialogue.

## Actual run and checks

Processed 50 TXT screenplay files representing all 50 episodes and generated 123 candidates. Candidate durations range from 25 to 60 seconds. 122 candidates contain dialogue evidence within the first eight seconds and 122 contain dialogue evidence within the final nine seconds. CSV, JSON and run summary are included.

## Scoring and assumptions

Scores are deterministic heuristics: opening dialogue bonus, ending dialogue bonus, bounded tension-term count, timed-beat density and small logline-overlap credit. Starts are selected from screenplay beat timestamps; durations are tested at 25, 30, 35, 40, 45, 50, 55 and 60 seconds. Overlap suppression rejects candidates whose shared interval is at least 25% of the shorter candidate.

## Limitations and failure cases

This is not a semantic AI editor. Keyword signals can misread context, sarcasm and character intent. Logline overlap is only a weak proxy for standalone clarity. A visual reveal may be stronger than dialogue, and a dialogue line may depend on earlier story context. Human editors must verify that the clip is understandable, emotionally compelling, visually coherent and ends on a genuine unresolved beat. No external LLM/API is called.

## Next steps and tools

The next iteration should add an LLM evaluator that receives each candidate plus surrounding context and returns structured, evidence-cited scores for hook strength, standalone clarity, emotional escalation and cliffhanger quality. Editor accept/reject labels can calibrate the ranking. Python and Streamlit implement the prototype; ChatGPT assisted with code drafting and review. A known weakness is that heuristic rules cannot reliably judge narrative meaning.